

CS 189/289


Today's lecture:

Linear regression (MLE + conditional Gaussians)

Assigned reading: 4-4.1.4


Regression

- *Supervised learning*: data pairs $D = \{(x_i, y_i)\}$, where, x_i may be discrete and/or continuous.
- *Regression*: label, y_i , is a real-valued, e.g., $y_i \in \mathbb{R}$.
- Formally, want $p(y|x)$, the conditional pdf.
- "Point" prediction is then $\hat{y} = E_Y[p(Y|X = x)]$.



Regression examples

- Covid infection rates from zip code and vaccination rate, etc.
- How much a particular protein will bind to a drug target.
- A person's blood pressure from their genetics.
- Tracking - object location in video at the next time-step.
- Housing prices, crime rates, stock prices, *etc.*
- Earliest regression: Legendre in 1805, and Gauss in 1809, both estimating orbits of bodies about the sun.




History of the term "Regression"

Sir Francis Galton (1822-1911) *"regression to the mean"*.



"It appeared from these experiments that the offspring did not tend to resemble their parents in size, but always to be more mediocre than they – to be smaller than the parents, if the parents were large; to be larger than the parents, if the parents were small."



History of the term "Regression"


Sir Francis Galton (1822-1911) *"regression to the mean"*.



"It appeared from these experiments that the


Don't confuse this meaning of the term "regression" with that in ML/statistics.

smaller than the parents, if the parents were large; to be larger than the parents, if the parents were small."



Possible Regression Tactics (goal: estimate $p(y|x)$)

- Our data are drawn from some distribution, $(X, Y) \sim p(x, y)$.
 - What are possible strategies to estimate $p(y|x)$?
1. Estimate $p(x, y|\theta)$ e.g. MVG for RVs X, Y , and then use fitted model to compute $p(y|x, \hat{\theta})$




Possible Regression Tactics (goal: estimate $p(y|x)$)

- Our data are drawn from some distribution, $(X, Y) \sim p(x, y)$.

- What are possible strategies to estimate $p(y|x)$?


1. Estimate $p(x, y|\theta)$ e.g. MVG for RVs X, Y , and then use fitted model to compute $p(y|x, \hat{\theta}) = \frac{p(y, x|\hat{\theta})}{p(x|\hat{\theta})} = \frac{p(y, x|\hat{\theta})}{\int_y p(y, x|\hat{\theta}) dy}$.



Possible Regression Tactics (goal: estimate $p(y|x)$)

- Our data are drawn from some distribution, $(X_i, Y_i) \sim p(x, y)$.
 - What are possible strategies to estimate $p(y|x)$?
1. Estimate $p(x, y|\theta)$ e.g. MVG for RVs X, Y , and then use fitted model to compute $p(y|x, \hat{\theta}) = \frac{p(y, x|\hat{\theta})}{p(x|\hat{\theta})} = \frac{p(y, x|\hat{\theta})}{\int_y p(y, x|\hat{\theta}) dy}$.
 2. Consider the inputs to be fixed, and model only the output as a RV. That is, directly model the conditional $p(y|x, \hat{\theta})$.

"generative", vs. "discriminative"



Linear Regression

- Takes the discriminative approach.
- Predictions are a linear function of the parameters:

$$\hat{y} = E_Y[p(y|x)] = \mathbf{w}^T \mathbf{x} + \mathbf{w}_0, \text{ for } \mathbf{w}, \mathbf{x} \in \mathbb{R}^d.$$


- \mathbf{w}_0 is called the "offset"/"bias"/"intercept".

Book-keeping trick: instead of a bias, we can make an extra feature that is always $\mathbf{1}$, now use $\mathbf{x}' = [\mathbf{x}, \mathbf{1}]$ and $\hat{y} = \mathbf{w}^T \mathbf{x}'$.

How useful can a linear model be?!

Which of these curves could have been modelled by linear regression?

$$\hat{y} = E_Y[p(y|x)] = \mathbf{w}^T \mathbf{x} + \mathbf{w}_0, \text{ for } \mathbf{w}, \mathbf{x} \in \mathbb{R}^d$$




How useful can a linear model be?!

$$w, x \in \mathbb{R}^1$$


$$\hat{y} = w^T x$$

Linear




$$\hat{y} = w^T [x, x^2]$$

Quadratic



$$\hat{y} = w^T [x, x^2, x^3]$$

Cubic



For full generality, $x \in \mathbb{R}^D$ need the cross-terms and bias terms for arbitrary polynomial, e.g., quadratic $[1, x_1, x_2, x_1^2, x_2^2, x_1 x_2]$.

Basis expansion of raw input space

$$x \in \mathbb{R}^{d=2} = [x_1, x_2] \rightarrow [1, x_1, x_2, x_1x_2, x_1^2, x_2^2] \in \mathbb{R}^{k=6}$$

Polynomial expansion of order 2 (i.e. quadratic)

- Denote *basis expansion* of the (raw) input features:
 $\Phi(x): \mathbb{R}^d \rightarrow \mathbb{R}^k$.

For $d = 1$, some polynomial expansions are:

- A quadratic expansion ($k = 2$), $\Phi(x) = [1, x, x^2]$.
- A cubic expansion ($k = 3$), $\Phi(x) = [1, x, x^2, x^3]$.

Identity basis expansion, $\Phi(x) = x$, and $k = d$.

Basis expansion of raw input space

Basis functions are pre-determined, so just a notational change:

$$\hat{y} = E_Y[p(y|x)] = \mathbf{w}^T \Phi(x), \text{ for } \mathbf{w} \in \mathbb{R}^k, x \in \mathbb{R}^d$$


In this lecture, for simplicity of notation, we will assume that this expansion has already been done, and just write $\hat{y} = \mathbf{w}^T \mathbf{x}$.

Many basis possible functions!

Polynomials: $f(x) = w_1 + w_2 x(1) + w_3 x(2) + w_4 x(1)^2 + w_5 x(2)^2 + w_6 x(1)x(2)$

$\phi(x) = (1, x(1), x(2), x(1)^2, x(2)^2, x(1)x(2))$.

Radial Basis fns:




$$\exp\left(\frac{-(x-c)^2}{r^2}\right)$$

Fourier basis:




Wavelets:




Specific form of linear regression

- So far we said linear regression is $\hat{y} = E_Y[p(y|x)] = E[p(y|x)] = \mathbf{w}^T \mathbf{x}$.
- But what do we use for $p(y|x)$?
- Standard linear regression uses a Gaussian $p(y|x) = N(y|\mathbf{w}^T \mathbf{x}, \sigma^2)$.
- Equivalent to $Y = \mathbf{w}^T \mathbf{x} + \epsilon$, with $\epsilon \sim N(0, \sigma^2)$.
- Which is equivalent to $Y - \mathbf{w}^T \mathbf{x} = \epsilon \sim N(0, \sigma^2)$.
- Alternate forms give "heavier tails" to the distribution of the "residual".



Aside: heavy-tailed distribution


- More of the mass lies further from the center of mass.
- Heavy-tailed noise better models outliers than a Gaussian.



$$Y - w^T x = \epsilon \sim N(0, \sigma^2)$$
$$Y - w^T x = \epsilon \sim \text{Cauchy}(0, \sigma^2)$$

Gaussian linear regression, $p(y|x) = N(y|\mathbf{w}^T \mathbf{x}, \sigma^2)$

For every value $\mathbf{X} = \mathbf{x}$, the target variable, Y , takes on a Gaussian distribution with the same variance, σ^2 :



"Training" a Gaussian linear regression model

How will we fit the regression model, $p(y|x) = N(y|\mathbf{w}^T x, \sigma^2)$?

$$\begin{aligned}\text{MLE: } \theta_{MLE} &= (w_{MLE}, \sigma_{MLE}^2) = \arg \max_{(w, \sigma^2)} \log p(D = \{(x_i, y_i)_{i=1}^n\} | \theta) \\ &= \arg \max_{(w, \sigma^2)} \sum_{i=1}^n \log p(y_i | x_i, \theta) \\ &= \arg \max_{(w, \sigma^2)} \sum_{i=1}^n \log N(y_i | w^T x_i, \sigma^2) \\ &= \arg \max_{(w, \sigma^2)} \sum_{i=1}^n \log \left(\frac{1}{\sqrt{2\pi\sigma^2}} \right) - \frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - w^T x_i)^2 \\ &= \arg \max_{(w, \sigma^2)} n \log \left(\frac{1}{\sqrt{2\pi\sigma^2}} \right) - \frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - w^T x_i)^2\end{aligned}$$

"Training" a Gaussian linear regression model

$$p(y|x) = N(y|w^T x, \sigma^2)$$

Note, if we ignore σ_{MLE}^2

$$(w_{MLE}, \sigma_{MLE}^2) = \arg \max_{(w, \sigma^2)} n \log \left(\frac{1}{\sqrt{2\pi\sigma^2}} \right) - \frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - w^T x_i)^2$$

... then estimating w above is the same as

$$\begin{aligned} &= \arg \min_w \sum_{i=1}^n (y_i - w^T x_i)^2 \\ &= \arg \min_w \sum_{i=1}^n (y_i - \hat{y}_i)^2 \end{aligned}$$

"least squares" loss function!

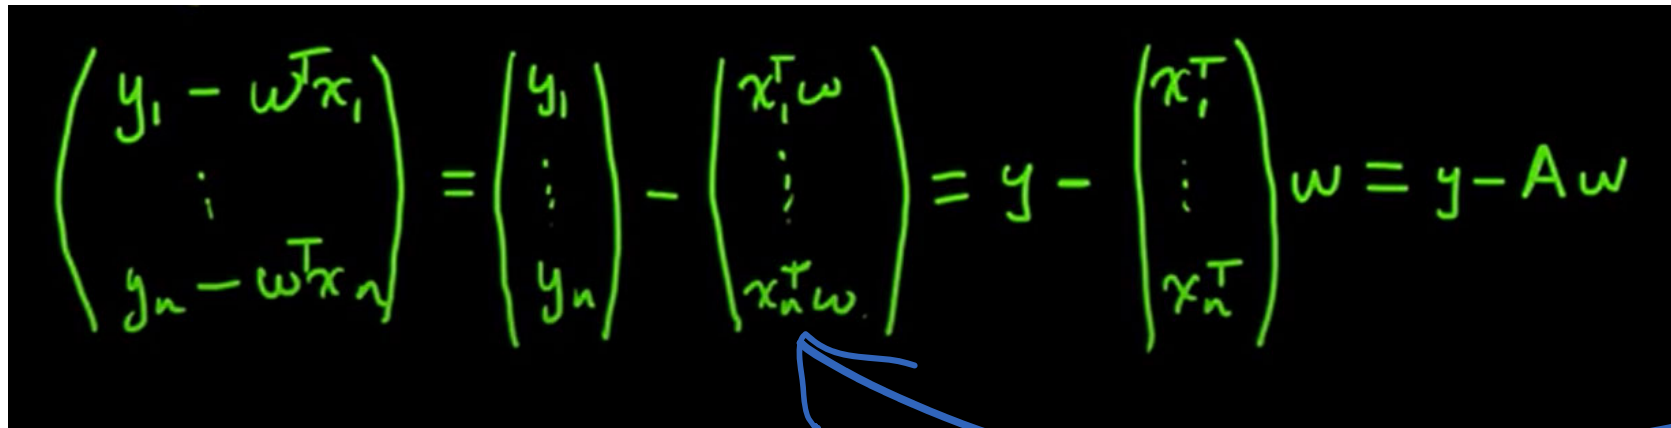
i.e. least squares arises naturally from conditional Gaussian MLE

"Training" a Gaussian linear regression model

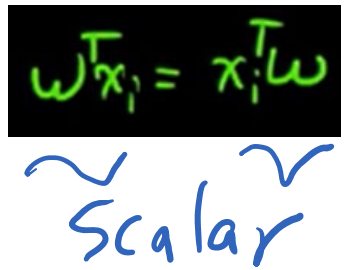
$$w_{\text{MLE}} = \arg \min_w \sum_{i=1}^n (y_i - w^T x_i)^2$$

Lets re-write this loss in "vectorized" form:

First, define:



A handwritten derivation on a black background showing the vectorization of the loss function. The equation is:
$$\begin{pmatrix} y_1 - w^T x_1 \\ \vdots \\ y_n - w^T x_n \end{pmatrix} = \begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix} - \begin{pmatrix} x_1^T w \\ \vdots \\ x_n^T w \end{pmatrix} = y - \begin{pmatrix} x_1^T \\ \vdots \\ x_n^T \end{pmatrix} w = y - A w$$



A handwritten note on a black background stating:
$$w^T x_i = x_i^T w$$
 Below the equation, the word "Scalar" is written with a checkmark above it. A blue arrow points from this note to the term $x_i^T w$ in the main derivation.

"Training" a Gaussian linear regression model

$$w_{\text{MLE}} = \arg \min_w \sum_{i=1}^n (y_i - w^T x_i)^2$$

$$A \in \mathbb{R}^{n \times d}$$

Lets re-write this loss in "vectorized" form:

First, define:

$$\begin{pmatrix} y_1 - w^T x_1 \\ \vdots \\ y_n - w^T x_n \end{pmatrix} = \begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix} - \begin{pmatrix} x_1^T w \\ \vdots \\ x_n^T w \end{pmatrix} = y - \begin{pmatrix} x_1^T \\ \vdots \\ x_n^T \end{pmatrix} w = y - A w$$

$$w^T x_i = x_i^T w$$

Scalar

A is called
the "design
matrix"

"Training" a Gaussian linear regression model

$$w_{\text{MLE}} = \arg \min_w \sum_{i=1}^n (y_i - w^T x_i)^2$$

$$A \in \mathbb{R}^{n \times d}$$

Lets re-write this loss in "vectorized" form:

First, define:

$$\begin{pmatrix} y_1 - w^T x_1 \\ \vdots \\ y_n - w^T x_n \end{pmatrix} = \begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix} - \begin{pmatrix} x_1^T w \\ \vdots \\ x_n^T w \end{pmatrix} = y - \begin{pmatrix} x_1^T \\ \vdots \\ x_n^T \end{pmatrix} w = y - Aw$$

$$w^T x_i = x_i^T w$$

Scalar

Then, we can re-write the loss as

$$\begin{aligned} & \arg \min_w (y - Aw)^T (y - Aw) \\ &= \arg \min_w \|y - Aw\|_2^2 \end{aligned}$$

A is called
the "design
matrix"

"Training" a Gaussian linear regression model

So want to minimize

$$\begin{aligned}\mathcal{L} &= \|y - Aw\|_2^2 & (y \in \mathbb{R}^n, A \in \mathbb{R}^{n \times d}, w \in \mathbb{R}^{d \times 1}) \\ &= (y - Aw)^T (y - Aw) \\ &= (y^T - (Aw)^T)(y - Aw) \\ &= y^T y - w^T A^T y - y^T A w + w^T A^T A w \\ &= y^T y - 2w^T A^T y + w^T A^T A w\end{aligned}$$

$$\begin{pmatrix} x_1^T \\ \vdots \\ x_n^T \end{pmatrix}$$

A

To minimize, we want to set $\frac{\partial \mathcal{L}}{\partial w} = 0$.

Most easily achieved by using the rules of vector calculus, so lets do a quick refresher.

$$p(y|x) = N(y|w^T x, \sigma^2)$$

Refresher on vector calculus

Some “rules” for taking gradients with respect to vectors.

- e.g., for vectors $\mathbf{a}, \mathbf{b} \in \mathbb{R}^{d \times 1}$, such that $\mathbf{a}^T \mathbf{b} \in \mathbb{R}$,

$$\frac{\partial(\mathbf{a}^T \mathbf{b})}{\partial a_j} = \frac{\partial(a_1 b_1 + a_2 b_2 + \cdots a_d b_d)}{\partial a_j} = b_j \in \mathbb{R}$$

Thus,

$$\frac{\partial(\mathbf{a}^T \mathbf{b})}{\partial \mathbf{a}} = \frac{\partial(a_1 b_1 + a_2 b_2 + \cdots a_d b_d)}{\partial \mathbf{a}} = \mathbf{b} \in \mathbb{R}^{d \times 1}$$

(not true for $\mathbf{a} \mathbf{b}^T$ which is a matrix, be careful!)

$$= \frac{\partial(\mathbf{b}^T \mathbf{a})}{\partial \mathbf{a}}$$

Refresher on vector calculus

For vector $\mathbf{x} \in \mathbb{R}^{d \times 1}$, and matrix $\Sigma \in \mathbb{R}^{d \times d}$

$$\frac{\partial \mathbf{x}^T \Sigma \mathbf{x}}{\partial \mathbf{x}} = (\Sigma + \Sigma^T) \mathbf{x}$$

Thus, if Σ is symmetric such that $\Sigma = \Sigma^T$ then

$$\frac{\partial \mathbf{x}^T \Sigma \mathbf{x}}{\partial \mathbf{x}} = 2\Sigma \mathbf{x}$$

(similar to the scalar version: $\frac{\partial (ax^2)}{\partial x} = 2ax$)

"Training" a Gaussian linear regression model

So want to minimize

$$\mathcal{L} = (y - Aw)^T (y - Aw)$$


$$(y \in \mathbb{R}^n, A \in \mathbb{R}^{n \times d}, w \in \mathbb{R}^{d \times 1})$$

$$= (y^T - (Aw)^T)(y - Aw)$$

$$= y^T y - w^T A^T y - y^T A w + w^T A^T A w$$

$$= y^T y - 2w^T A^T y + w^T A^T A w$$

To minimize, we want to set $\frac{\partial \mathcal{L}}{\partial w} = 0$.



$$\begin{pmatrix} x_1^T \\ \vdots \\ x_n^T \end{pmatrix} \Rightarrow A \in \mathbb{R}^{n \times d}$$

"Training" a Gaussian linear regression model

So want to minimize

$$\mathcal{L} = (y - Aw)^T (y - Aw)$$

$$(y \in \mathbb{R}^n, A \in \mathbb{R}^{n \times d}, w \in \mathbb{R}^{d \times 1})$$


$$\begin{pmatrix} x_1^T \\ \vdots \\ x_n^T \end{pmatrix} = A \in \mathbb{R}^{n \times d}$$

$$= (y^T - (Aw)^T)(y - Aw)$$

$$= y^T y - w^T A^T y - y^T A w + w^T A^T A w$$

$$= y^T y - 2w^T A^T y + w^T A^T A w$$

To minimize, we want to set $\frac{\partial \mathcal{L}}{\partial w} = 0$.

$$\nabla_w \mathcal{L} = \begin{bmatrix} \frac{\partial \mathcal{L}}{\partial w_1} \\ \frac{\partial \mathcal{L}}{\partial w_2} \\ \vdots \\ \frac{\partial \mathcal{L}}{\partial w_d} \end{bmatrix} = \begin{bmatrix} -2A^T y + 2A^T A w \end{bmatrix}$$
$$= \underbrace{(-A^T y + A^T A w)}_{d \times 1} \cdot 2$$

"Training" a Gaussian linear regression model

So want to minimize

$$\begin{aligned}\mathcal{L} &= (y - Aw)^T (y - Aw) \\ &= (y^T - (Aw)^T)(y - Aw) \\ &= y^T y - w^T A^T y - y^T A w + w^T A^T A w \\ &= y^T y - 2w^T A^T y + w^T A^T A w\end{aligned}$$

To minimize, we want to set $\frac{\partial \mathcal{L}}{\partial w} = 0$.

$$\nabla_w \mathcal{L} = \begin{bmatrix} \frac{\partial \mathcal{L}}{\partial w_1} \\ \frac{\partial \mathcal{L}}{\partial w_2} \\ \vdots \\ \frac{\partial \mathcal{L}}{\partial w_n} \end{bmatrix} = \begin{bmatrix} -2A^T y + 2A^T A w \end{bmatrix} = \underbrace{-A^T y + A^T A w}_{d \times 1} \cdot 2$$

"left pseudo-inverse" of A
 $(A^T A)^{-1} A^T * A = I$

$$A^+ A = I$$

$$\begin{pmatrix} x_1^T \\ \vdots \\ x_n^T \end{pmatrix} = A \in \mathbb{R}^{n \times d}$$

set to 0

rectangular
so no A^{-1} !

$$-A^T y + A^T A w = 0$$

$$A^T A w = A^T y$$

$$w = (A^T A)^{-1} A^T y$$

"Training" a Gaussian linear regression model

- A^+ is the Moore-Penrose inverse.
- Can be used even when A is not full rank.
- i.e., when A has dependent feature vectors.
- Will yield w_{MLE} with min. 2-norm, $\|w_{MLE}\|_2$

Related to spectral decomposition from last class:

Inverses and square roots

- If $A = \Phi D \Phi^T$
Then $A^{-1} = \Phi D^{-1} \Phi^T$
where $D^{-1} = \begin{bmatrix} \lambda_1^{-1} & & \\ & \lambda_2^{-1} & \\ & & \ddots \\ & & & \lambda_n^{-1} \end{bmatrix}$

Take reciprocal
of only non-
zero values,
leave the rest
as zeros.

"left pseudo-inverse" of A
 $(A^T A)^{-1} A^T A = I$

$$A^+ A = I$$

$$\begin{pmatrix} x_1^T \\ \vdots \\ x_n^T \end{pmatrix} = A \in \mathbb{R}^{n \times d}$$

rectangular
So no A^{-1} !

$$-A^T y + A^T A w = 0$$

$$A^T A w = A^T y$$

$$w = (A^T A)^{-1} A^T y$$

"Training" a Gaussian linear regression model

- We still need to check if the critical point, $(A^T A)^{-1} A^T y$ is minimum of the squared error loss.

$$w =$$

$$\begin{pmatrix} x_1^T \\ \vdots \\ x_n^T \end{pmatrix} = A \in \mathbb{R}^{n \times d}$$

- Recall $\nabla_w \mathcal{L} = -2A^T y + 2A^T A w$
- So Hessian matrix ($\nabla_w^2 \mathcal{L}$) is $2A^T A$. When is $A^T A$ PD?
- When the features in data set are independent (when it has full rank).

σ^2 from MLE is just the mean squared residual, $\sigma^2 = \frac{1}{N} \sum_i (y_i - w^T x)^2$.

$$\nabla_w \mathcal{L} = \begin{bmatrix} \partial \mathcal{L} / \partial w_1 \\ \partial \mathcal{L} / \partial w_2 \\ \vdots \\ \partial \mathcal{L} / \partial w_d \end{bmatrix} = \begin{bmatrix} -2A^T y + 2A^T A w \end{bmatrix}$$
$$= \underbrace{-A^T y + A^T A w}_{d \times 1}$$


$$-A^T y + A^T A w = 0$$

$$A^T A w = A^T y$$

$$\boxed{w = (A^T A)^{-1} A^T y}$$

Regression in 1950s

Electromechanical desk "calculators" were used, and it could take up to 24 hours to receive the result from one regression on a small data set.




Geometric view of linear regression

- We're trying to predict all our training data labels correctly, such that $y_i = w^T x_i$ for all $i \in [1 \dots n]$.
- In vector form, this means we're looking for

$$\begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix} = A w = \begin{bmatrix} x_1^T \\ x_2^T \\ \vdots \\ x_n^T \end{bmatrix} w$$


Generally, this is not possible because of noise; or incorrect model (e.g. missing some features).

$$\begin{pmatrix} x_1^T \\ \vdots \\ x_n^T \end{pmatrix} = A$$




Geometric view of linear regression, $y = Aw$

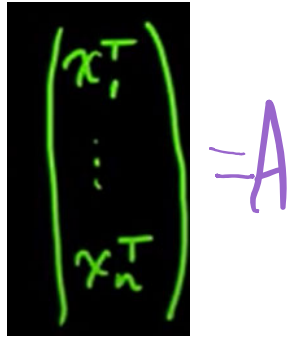
- So let's think about the **error vector** ($\mathbf{e} \equiv \mathbf{y} - \hat{\mathbf{y}} = \mathbf{y} - A\mathbf{w}, \in \mathbb{R}^{n \times 1}$).
- A "good" setting of \mathbf{w} minimizes the magnitude of \mathbf{e} .
- Magnitude is minimized when \mathbf{e} lies \perp to column space of A .
- Thus we seek \mathbf{w} such that $\mathbf{e}^T A = A^T \mathbf{e} = \mathbf{0} = A^T(\mathbf{y} - A\mathbf{w})$.
- Thus $A^T \mathbf{y} - A^T A \mathbf{w} = \mathbf{0}$.
- Thus $A^T \mathbf{y} = A^T A \mathbf{w}$ (same as from MLE/least squares)!
- Thus $\mathbf{w} = (A^T A)^{-1} A^T \mathbf{y}$, as before.



$$\begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{pmatrix} = A\mathbf{w} = \begin{bmatrix} x_1^T \\ x_2^T \\ \vdots \\ x_N^T \end{bmatrix} \mathbf{w}$$



Using basis expansions instead of x_i


$$\begin{pmatrix} x_1^T \\ \vdots \\ x_n^T \end{pmatrix} = A$$

- $\{x_j\} \rightarrow \{\Phi(x_j)\}$?
- Just define \mathbf{A} with $\Phi^T(x_j)$ instead of x_j because $\Phi(x)$ is fixed ahead of time, so its like someone just gave us different raw inputs x .




$$x = \Phi[x_1, x_2] = [1, x_1, x_2, x_1 x_2, x_1^2, x_2^2] \in \mathbb{R}$$

Polynomial expansion of order 2 (i.e. quadratic)

What can go wrong in linear regression

As we add higher and higher order polynomials, a few things happen:

1. # features, d gets bigger and bigger.
2. For $d \geq n$ can perfectly fit any data (i.e., polynomials are a complete basis).
3. Even when we don't perfectly fit the training data, we are still in danger of overfitting (worse prediction on test set). Our goal is not to fit a line through the training data exactly, it is to do well on unseen test cases!



What can go wrong in linear regression

Two main categories of fixes:

1. Remove features until the problem is well-behaved (“feature selection”).
2. Leave the features as they are, but *add constraints to the system* to “tighten it up” (aka “regularization”)—next class.

NB: Moore-Penrose inverse is not a general fix for ML models, only works for linear regression.

A thought experiment

Consider:

- Use MLE on data, $D = \{x_i, y_i\}$ to get $\hat{y}_i = p_{\theta}(y|\Phi(x))$ using linear regression.
- Assume an abundance of data (millions of data points), and only 100 parameters.
- Suppose get accuracy $\mp \$1000$ of sale price when applying to held out part of our data.
- Can we assume this model will get $\mp \$1000$ on any test set that may come in the future?

Actual vs. predicted sale price of house




Fig. 4 Ridge Prediction for Training Data.

Causation vs correlation

Breakingviews

Zillow's failed house flipping

Reuters

WSJ NOV. 2021 : “*The company expects to record losses of more than \$500 million from home-flipping by the end of this year and is laying off a quarter of its staff.*”

Actual vs. predicted sale price of house




Fig. 4 Ridge Prediction for Training Data.