## Gradient Descent and Backpropagation 1

Saeed Saremi

Assigned reading: 6.{2,4}, 7.{1,2}, 5.4.4

September 24, 2024

### a summary of the previous lecture

• the geometry of  $\theta$  in the logistic regression given the following parametrization:<sup>1</sup>

$$p(1|x) = \frac{1}{1 + \exp(-\theta^{\top}x)} =: p(z(x;\theta)), \text{ where } z(x;\theta) = \theta^{\top}x$$

 $\theta$  is perpendicular to the decision boundary and points to class "1".

- softmax function: generalization of the logistic regression to multiclass (K > 2) classification via the generative approach, where we assumed  $X | k \sim \mathcal{N}(\mu_k, \Sigma)$ .
- ▶ the negative log-likelihood loss for the logisitic regression and its gradient:

$$\mathcal{L}(\theta) = -y \log p(z(x;\theta)) - (1-y) \log(1 - p(z(x;\theta)))$$
$$\nabla_{\theta} \mathcal{L}(\theta) = (p_z - y) \nabla_{\theta} z = (p_z - y) \times,$$

and we interpreted this gradient geometrically in terms of the parameter updates.

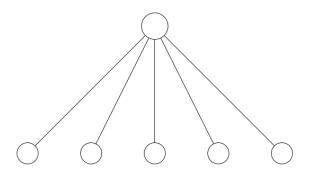
▶ artificial neural networks as simplified models of biological neural networks

<sup>&</sup>lt;sup>1</sup>Here, I assume the data is "centered", therefore  $\theta_0 = 0$ .

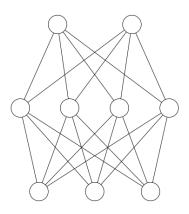
#### outline

- neural networks as a certain type of nonlinear functions
- > (linear/logisitic regression as single-layer neural network)
- > finish the comparisons with biological neural networks
- training neural networks: stochastic gradient descent
- > gradient of a function
- > the chain rule
- > convex and non-convex functions
- > geometrical interpretation of the gradient
- > stochastic gradient descent (SGD)
- **>** a prelude to backpropagation: the cross-entropy loss for multiclass (K > 2) classification and its gradient

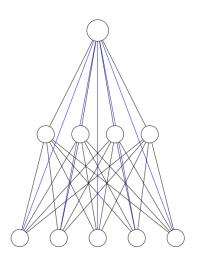
# single-layer neural networks

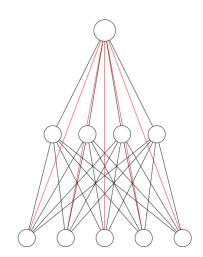


# two-layer neural networks



complex neural network architectures:anything goes, but "loops" are not allowed!





## biological neural networks v.s. deep neural networks



Figure: human brain:  $10^{11}$  neurons,  $10^{15}$  synapses

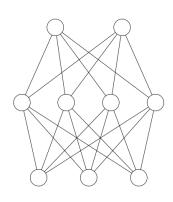
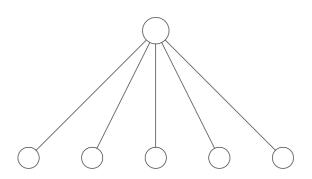


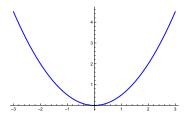
Figure: GPT-4: 10<sup>6</sup> neurons, 10<sup>11</sup> parameters (weights), 100 layers

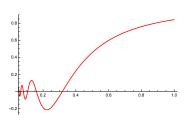
**?** What are we missing in this comparison?

# single-layer neural networks: the linear and logistic regression



#### convex and non-convex functions





Recall from the multivariate calculus that the gradient of the function  $f: \mathbb{R}^m \to \mathbb{R}$  is the vector

$$\nabla f(\theta_1, \dots, \theta_m) = \left(\frac{\partial f}{\partial \theta_1}, \dots, \frac{\partial f}{\partial \theta_m}\right)^{\top}$$

- ▶ The gradient is the direction that leads to maximal increase of *f* (steepest ascent). Equivalently, the negative gradient is the direction of steepest descent.
- ▶ Formally, a function f is convex if for all  $\theta_1$ ,  $\theta_2$  in  $\mathbb{R}^m$  and  $t \in [0,1]$ :

$$f(t\theta_1 + (1-t)\theta_2) \le f(\theta_1) + (1-t)f(\theta_2)$$

• convex functions have a single (global) minimum.

## review: the chain rule from multivariate calculus

Given  $u(z_1, z_2) = (z_1 + z_2)^2/2$ , where  $z_1(x_1, x_2) = x_1 \sin x_2$  and  $z_2(x_1, x_2) = (\sin x_2)^2$ ; determine  $\nabla_x u$  using the chain rule:

brute force:

chain rule:

## the geometrical meaning of the gradient

Prove that  $\nabla f(\theta)$ , the gradient of the function  $f:\Theta\to\mathbb{R}$  at any point  $\theta$ , is perpendicular to the level set of f at that point.<sup>2</sup>

(Hint: define a path  $t \to \theta$  on the level set and use the fact that by definition of the level set  $\partial_t f(\theta(t)) = 0$ . Use the chain rule!)

<sup>&</sup>lt;sup>2</sup>Recall that level sets are defined by  $\{\theta': f(\theta') = f(\theta)\}$  for some constant c.

## stochastic gradient descent I

▶ at a high level the loss is always written as the sum of losses by individual points  $i \in [n] := \{1, ..., n\}$  in the training set  $\mathcal{D} = \{(x_i, y_i)\}_{i \in [n]}$ :

$$\mathcal{L}(\theta) = \sum_{i=1}^{n} \mathcal{L}_{i}(\theta),$$

where  $\mathcal{L}_i(\theta)$  is short for:

$$\mathcal{L}_i(\theta) := \mathcal{L}(x_i, y_i; \theta).$$

▶ This is very general, but it's very easy to see where it's coming from in the maximum log-likelihood framework. We always assume the i.i.d. setting:

$$(x_i, y_i) \stackrel{\text{iid}}{\sim} p_{\theta}(x, y), i \in [n].$$

Therefore,

$$p(\mathcal{D}|\theta) = \prod_{i=1}^n p_{\theta}(x_i, y_i).$$

It follows:

$$\mathcal{L}_i(\theta) = -\log p_{\theta}(x_i, y_i).$$

## stochastic gradient descent II

The two prominent examples so far include:

▶ linear regression:

$$-\log p_{\theta}(x_i, y_i) = \frac{1}{2\sigma^2}(y_i - \theta^{\top} x_i^{(0)}) + const$$

▶ logistic "regression":

$$-\log p_{\theta}(x_{i}, y_{i}) = -y_{i} \log g(\theta^{\top} x_{i}^{(0)}) - (1 - y_{i}) \log(1 - g(\theta^{\top} x_{i}^{(0)})) + const$$

Note that the structure of the loss is general: we are getting ready to replace  $\theta^{\top} x_i^{(0)}$  with an *L*-layer neural network:  $f(x_i^{(0)}; \theta^{(1)}, \dots, \theta^{(L)})$ .)

Stochastic Gradient Descent (SGD) in its pure form is defined by the following updates:

$$\theta_{t+1} = \theta_t - \epsilon_t \nabla_{\theta} \mathcal{L}(\mathbf{x}_i, \mathbf{y}_i; \theta),$$

where  $i \in [n]$  is selected at random (typically without replacement) at each iteration t.

## stochastic gradient descent III

- ▶ Intuitively (people have tried to study this) in high dimensions the loss landscape is "dominated" by saddle points and the noise in SGD helps to avoid them.
- ▶ SGD is convenient in the regime  $n \gg 1$ .
- One pass through the data is called an epoch.
- ▶ The problem is towards the end of the training (optimization) the noise in SGD will slow down the training.
- ▶ In short: noise helps us at the beginning of training, it "hurts" us towards the end.
- ▶ Of course, one can find a compromise by dividing the dataset into (random) mini-batches of size b: in this scheme one epoch involves  $\lfloor n/b \rfloor$  updates.
- ? Based on this picture, can you suggest a batching scheme for effective training? <sup>3</sup>

 $<sup>^3</sup>$ Coming up with a mini-batch / learning rate schedule remains an art and problem dependent.

# cross-entropy loss for multiclass classification and its gradient