

Backpropagation and SGD

Saeed Saremi

Assigned reading: 5.4.4, 8

September 26, 2024

a summary of the previous lecture(s) I

- ▶ L -layer (feedforward) **neural network** as a particular way of composing functions:

$$\begin{aligned} a_j^{(\ell)} &= \sum_{i \in [d_{\ell-1}]} \theta_{ji}^{(\ell)} x_i^{(\ell-1)} \\ x_j^{(\ell)} &= g_\ell(a_j^{(\ell)}), \\ g_\ell &: \mathbb{R} \rightarrow \mathbb{R}, \end{aligned}$$

where the computation is iterative from $\ell = 1$ to $\ell = L$.

- ▶ $x^{(0)}$ represents the INPUT, $x^{(L)}$ the OUTPUT: $f_\theta : x^{(0)} \mapsto x^{(L)}$
- ▶ by convention $g_L(z) = z$, and $g_\ell = g$ is typically fixed for $\ell \in [L-1]$.
- ▶ examples of g include: $z \mapsto 1/(1 + e^{-z})$ and $z \mapsto \max(z, 0)$.
- ▶ one can increase the “**capacity**” of the neural network by either increasing L or increasing the number of **hidden neurons** (and thus connections) in each layer.

a summary of the previous lecture(s) II

- ▶ The loss $\mathcal{L}(\theta)$ we minimize is formulated as negative log likelihood.

> regression, $\mathbb{R}^d \rightarrow \mathbb{R}$:

$$\mathcal{L} = \frac{1}{2}(y - a_1^{(L)})^2$$

> classification, $\mathbb{R}^d \rightarrow \{0, 1\}$:

$$\mathcal{L} = -y \log \sigma(a_1^{(L)}) - (1 - y) \log(1 - \sigma(a_1^{(L)})),$$

where σ is the logistic sigmoid function:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- ▶ for a 1-layer ($L = 1$) neural network:

$$a_k^{(L)}(x^{(0)}; \theta) = \sum_{i \in [d_0]} \theta_{ki}^{(1)} x_i^{(0)},$$

where in the two examples above $k = 1$ (only one output neuron).

a summary of the previous lecture(s) III

- ▶ In **gradient**-based optimization we minimize the loss $\mathcal{L}(\theta)$ using **gradient descent**:

$$\theta_{t+1} = \theta_t - \epsilon_t \nabla_{\theta_t} \mathcal{L}(\theta_t),$$

where ϵ_t is some “small” number where we can take to be constant $\epsilon_t = \epsilon$.

- ▶ For 1-layer ($L = 1$) neural networks $\nabla \mathcal{L}$ takes the following general form:

$$\partial_{\theta_{ki}^{(1)}} \mathcal{L} = \partial_{a_k^{(L)}} \mathcal{L} \partial_{\theta_{ki}^{(1)}} a_k^{(L)}$$

- ▶ (gaussian) regression:

$$\partial_{a_1^{(L)}} \mathcal{L} = a_1^{(L)} - y$$


- ▶ logistic regression:

$$\partial_{a_1^{(L)}} \mathcal{L} = \sigma(a_1^L) - y$$

outline I

- ▶ $\partial_{a^{(L)}} \mathcal{L}$ has a clear interpretation in terms of the error signal that drives **learning**.
- ▶ We therefore name it as the **error** (signal) and use the following notation:


$$\Delta_k^{(L)} = \partial_{a_k^{(L)}} \mathcal{L}.$$

-  Next we derive $\Delta_k^{(L)}$ for multi-class ($K > 2$) classification, where the loss is given by

$$\mathcal{L} = - \sum_{k \in [K]} y_k \log \frac{\exp(a_k^{(L)})}{Z},$$

where Z is the normalization:

$$Z = \sum_{k \in [K]} \exp(a_k^{(L)})$$

-  Can you guess the answer?

outline II

- ✎ Derive an algorithm for computing

$$\partial_{\theta_{ji}^{(l)}} \mathcal{L}$$

for **all** the elements of $\theta = (\theta^{(1)}, \dots, \theta^{(L)})$ that is of $O(m)^1$ where $m = \dim(\theta)$.

- The **error** signals

$$\Delta_j^{(l)} = \frac{\partial \mathcal{L}}{\partial a_j^{(l)}}$$

plays a central role in the final algorithm!

- The algorithm called **backpropagation** involves computing the errors backward from top $\ell = L$ back to $\ell = 1$ by a **single pass** through the neural network.²

¹We already know a (**doubly bad**) algorithm that is of $O(m^2)$:

$$\partial_{\theta_{ji}^{(l)}} \mathcal{L} \approx \frac{\mathcal{L}(\theta_{ji}^{(l)} + \epsilon) - \mathcal{L}(\theta_{ji}^{(l)} - \epsilon)}{2\epsilon}$$

²I recommend reading the original paper: D. E. Rumelhart, G. E. Hinton, and R. J. Williams, *Learning representations by back-propagating errors*, Nature, 1986.

stochastic gradient descent I

- ▶ at a high level the loss is **always** written as the **sum** of losses by individual points $i \in [n] := \{1, \dots, n\}$ in the training set $\mathcal{D} = \{(x_i, y_i)\}_{i \in [n]}$:

$$\mathcal{L}(\theta) = \sum_{i=1}^n \mathcal{L}_i(\theta),$$

where $\mathcal{L}_i(\theta)$ is short for:

$$\mathcal{L}_i(\theta) := \mathcal{L}(x_i, y_i; \theta).$$

- ▶ This is very general, but it's very easy to see where it's coming from in the maximum **log-likelihood** framework. We **always** assume the i.i.d. setting:

$$(x_i, y_i) \stackrel{\text{iid}}{\sim} p_{\theta}(x, y), \quad i \in [n].$$

Therefore,

$$p(\mathcal{D}|\theta) = \prod_{i=1}^n p_{\theta}(x_i, y_i).$$

It follows:

$$\mathcal{L}_i(\theta) = -\log p_{\theta}(x_i, y_i).$$

stochastic gradient descent II

Stochastic Gradient Descent (SGD) in its pure form is defined by the following updates:

$$\theta_{t+1} = \theta_t - \epsilon_t \nabla_{\theta_t} \mathcal{L}_i(\theta_t),$$

where $i \in [n]$ is selected at random (typically without replacement) at each iteration t .

stochastic gradient descent III

- ▶ Intuitively (people have tried to study this) in **high dimensions** the loss landscape is “dominated” by **saddle points** and the **noise** in **SGD** helps to avoid them.
- ▶ SGD is **convenient** in the regime $n \ggg 1$.
- ▶ One pass through the data is called an **epoch**.
- ▶ The problem is towards the end of the training (optimization) the noise in SGD will slow down the training.
- ▶ In short: noise helps us at the beginning of training, it “hurts” us towards the end.
- ▶ Of course, one can find a compromise by dividing the dataset into (random) **mini-batches** of size b : in this scheme one **epoch** involves $\lfloor n/b \rfloor$ **updates**.
- ❓ Based on this picture, can you suggest a **batching scheme** for effective training?³

³Coming up with a mini-batch/learning rate schedule remains an art and it is problem dependent.